

Training Day-79 Report:

1. Sentiment Analysis Hands-On

Sentiment Analysis is a common NLP task where the goal is to determine the sentiment (e.g., positive, negative, neutral) of a given text.

Steps for Sentiment Analysis:

1. **Data Preparation:** Use a dataset like IMDb Movie Reviews or any labeled sentiment dataset.
2. **Preprocessing:**
 - Tokenization.
 - Padding sequences to a fixed length.
 - Encoding labels.
3. **Building the Model:** Use an Embedding layer followed by LSTM/GRU or CNN for feature extraction.
4. **Training:** Train the model to classify sentiments.
5. **Evaluation:** Test the model on unseen data.

Implementation in TensorFlow:

```
import tensorflow as tf

from tensorflow.keras.datasets import imdb

from tensorflow.keras.preprocessing.sequence import pad_sequences

from tensorflow.keras.models import Sequential

from tensorflow.keras.layers import Embedding, LSTM, Dense, Dropout


# Load IMDb dataset

vocab_size = 10000

max_len = 200

(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=vocab_size)


# Pad sequences to ensure uniform input length

x_train = pad_sequences(x_train, maxlen=max_len)
```

```

x_test = pad_sequences(x_test, maxlen=max_len)

# Build the sentiment analysis model
model = Sequential([
    Embedding(input_dim=vocab_size, output_dim=128, input_length=max_len), #
    Embedding layer
    LSTM(64, return_sequences=False), # LSTM layer
    Dropout(0.5), # Dropout for regularization
    Dense(1, activation='sigmoid') # Output layer for binary
    classification
])

# Compile the model
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

# Train the model
model.fit(x_train, y_train, epochs=3, batch_size=64, validation_data=(x_test, y_test))

# Evaluate the model
loss, accuracy = model.evaluate(x_test, y_test)
print(f"Test Accuracy: {accuracy}")

# Make predictions
sample_review = "The movie was fantastic and very engaging!"
encoded_review = [1] + [word for word in sample_review.split() if word in
imdb.get_word_index()]
padded_review = pad_sequences([encoded_review], maxlen=max_len)
prediction = model.predict(padded_review)
print(f"Sentiment: {'Positive' if prediction[0][0] > 0.5 else 'Negative'}")

```

2. Seq-to-Seq Model

Sequence-to-Sequence (Seq2Seq) models are used for tasks like machine translation, summarization, and chatbot applications. These models take a sequence as input and generate another sequence as output.

Seq-to-Seq Architecture:

1. Encoder:

- Encodes the input sequence into a fixed-length context vector.
- Typically uses RNNs, LSTMs, or GRUs.

2. Decoder:

- Takes the context vector as input and generates the output sequence.

3. Attention Mechanism:

- Enhances performance by allowing the decoder to focus on specific parts of the input sequence at each time step.

Seq-to-Seq Implementation for Machine Translation:

```
import tensorflow as tf

from tensorflow.keras.layers import Input, LSTM, Dense
from tensorflow.keras.models import Model

import numpy as np

# Sample data
input_texts = ["hello", "how are you", "good morning"]
target_texts = ["hola", "cómo estás", "buenos días"]

# Tokenize data
input_tokenizer = tf.keras.preprocessing.text.Tokenizer()
target_tokenizer = tf.keras.preprocessing.text.Tokenizer()

input_tokenizer.fit_on_texts(input_texts)
target_tokenizer.fit_on_texts(target_texts)

input_sequences = input_tokenizer.texts_to_sequences(input_texts)
```

```
target_sequences = target_tokenizer.texts_to_sequences(target_texts)
```

```
input_data = tf.keras.preprocessing.sequence.pad_sequences(input_sequences,  
padding='post')
```

```
target_data = tf.keras.preprocessing.sequence.pad_sequences(target_sequences,  
padding='post')
```

```
# Define model parameters
```

```
num_encoder_tokens = len(input_tokenizer.word_index) + 1
```

```
num_decoder_tokens = len(target_tokenizer.word_index) + 1
```

```
latent_dim = 256
```

```
# Encoder
```

```
encoder_inputs = Input(shape=(None,))
```

```
encoder_embedding = Dense(64, activation='relu')(encoder_inputs)
```

```
encoder_lstm = LSTM(latent_dim, return_state=True)
```

```
encoder_outputs, state_h, state_c = encoder_lstm(encoder_embedding)
```

```
# Decoder
```

```
decoder_inputs = Input(shape=(None,))
```

```
decoder_embedding = Dense(64, activation='relu')(decoder_inputs)
```

```
decoder_lstm = LSTM(latent_dim, return_sequences=True, return_state=True)
```

```
decoder_outputs, _, _ = decoder_lstm(decoder_embedding, initial_state=[state_h, state_c])
```

```
decoder_dense = Dense(num_decoder_tokens, activation='softmax')
```

```
decoder_outputs = decoder_dense(decoder_outputs)
```

```
# Define model
```

```
model = Model([encoder_inputs, decoder_inputs], decoder_outputs)
```

```
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy')
```

```
# Training
```

```
model.fit([input_data, target_data[:, :-1]], target_data[:, 1:], batch_size=64, epochs=50)
```

```
# Inference
```

```
def translate(input_seq):
```

```
    states = encoder_lstm.predict(input_seq)
```

```
    translated_seq = decoder_lstm.predict(states)
```

```
    return translated_seq
```

```
print("Translation:", translate(input_data[0:1]))
```

Summary:

1. Sentiment Analysis:

- Used LSTM for sequence-based binary classification.
- Applied text preprocessing, tokenization, and padding.

2. Seq-to-Seq:

- Demonstrated architecture for sequence translation tasks.
- Encoder-decoder structure handles input-to-output sequence transformation.